Y. Jiang, P. Foret, S. Yak, D. M. Roy, H. Mobahi, G. K. Dziugaite, S. Bengio, S. Gunasekar, I. Guyon, and B. Neyshabur. Neurips 2020 Competition: Predicting generalization in deep learning, 2020a. URL `https://arxiv.org/abs/2012.07976`.

Y. Jiang, B. Neyshabur, H. Mobahi, D. Krishnan, and S. Bengio. Fantastic generalization measures and where to find them. In *International Conference on Learning Representations*, 2020b. URL `https://openreview.net/forum?id=SJgIPJBFvH`.

Y. Jiang, P. Natekar, M. Sharma, S. K. Aithal, D. Kashyap, N. Subramanyam, C. Lassance, D. M. Roy, G. K. Dziugaite, S. Gunasekar, I. Guyon, P. Foret, S. Yak, H. Mobahi, B. Neyshabur, and S. Bengio. Methods and analysis of the first competition in predicting generalization of deep learning. In H. J. Escalante and K. Hofmann, editors, *Proceedings of the NeurIPS 2020 Competition and Demonstration Track*, volume 133 of *Proceedings of Machine Learning Research*, pages 170–190. PMLR, 06–12 Dec 2021. URL `https://proceedings.mlr.press/v133/jiang21a.html`.

M. Kang, W. J. Shim, M. Cho, and J. Park. Rebooting ACGAN: Auxiliary classifier GANs with stable training. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. W. Vaughan, editors, *Advances in Neural Information Processing Systems*, 2021. URL `https://openreview.net/forum?id=r7UC-b67YkO`.

T. Karras, T. Aila, S. Laine, and J. Lehtinen. Progressive growing of GANs for improved quality, stability, and variation. In *International Conference on Learning Representations*, 2018. URL `https://openreview.net/forum?id=Hk99zCeAb`.

T. Karras, S. Laine, M. Aittala, J. Hellsten, J. Lehtinen, and T. Aila. Analyzing and improving the image quality of StyleGAN. In *Proc. CVPR*, 2020.

M. Kelly, R. Longjohn, and K. Nottingham. The UCI Machine Learning Repository. `https://archive.ics.uci.edu`.

A. Khovanskiĭ. *Fewnomials*. Translations of Mathematical Monographs, 88. American Mathematical Society, Providence, 1991. ISBN 0821845470.

V. Khrulkov and I. Oseledets. Geometry score: A method for comparing generative adversarial networks. In J. Dy and A. Krause, editors, *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 2621–2629. PMLR, 10–15 Jul 2018. URL `https://proceedings.mlr.press/v80/khrulkov18a.html`.

P. J. Kim, Y. Jang, J. Kim, and J. Yoo. Topp&r: Robust support estimation approach for evaluating fidelity and diversity in generative models. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023. URL `https://openreview.net/forum?id=2gUCMr6fDY`.

D. P. Kingma and M. Welling. Auto-encoding variational Bayes. In Y. Bengio and Y. LeCun, editors, *2nd International Conference on Learning Representations, ICLR 2014, Banff, AB, Canada, April 14-16, 2014, Conference Track Proceedings*, 2014. URL `http://arxiv.org/abs/1312.6114`.

N. Klyuchnikov, I. Trofimov, E. Artemova, M. Salnikov, M. Fedorov, and E. Burnaev. NAS-Bench-NLP: Neural architecture search benchmark for natural language processing, 2020. URL https://arxiv.org/abs/2006.07116.

G. Kozma, Z. Lotker, and G. Stupp. The minimal spanning tree and the upper box dimension. *Proceedings of the American Mathematical Society*, 134(4):1183–1187, 2006. ISSN 00029939, 10886826.

A. Krizhevsky. Learning multiple layers of features from tiny images. 2009. URL http://www.cs.toronto.edu/~kriz/cifar.html.

A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks. In F. Pereira, C. Burges, L. Bottou, and K. Weinberger, editors, *Advances in Neural Information Processing Systems*, volume 25. Curran Associates, Inc., 2012. URL https://proceedings.neurips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf.

S. V. Kurochkin. Neural network with smooth activation functions and without bottlenecks is almost surely a Morse function. *Computational Mathematics and Mathematical Physics*, 61(7):1162–1168, Jul 2021. ISSN 1555-6662. doi: 10.1134/S0965542521070101. URL https://doi.org/10.1134/S0965542521070101.

T. Lacombe, Y. Ike, M. Carrière, F. Chazal, M. Glisse, and Y. Umeda. Topological uncertainty: Monitoring trained neural networks through persistence of activation graphs. In Z.-H. Zhou, editor, *Proceedings of the Thirtieth International Joint Conference on Artificial Intelligence, IJCAI-21*, pages 2666–2672. International Joint Conferences on Artificial Intelligence Organization, 8 2021. doi: 10.24963/ijcai.2021/367. URL https://doi.org/10.24963/ijcai.2021/367. Main Track.

Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998. doi: 10.1109/5.726791.

J. Leygonie, S. Oudot, and U. Tillmann. A framework for differential calculus on persistence barcodes. *Foundations of Computational Mathematics*, 22(4):1069–1131, Aug 2022. ISSN 1615-3383. doi: 10.1007/s10208-021-09522-y. URL https://doi.org/10.1007/s10208-021-09522-y.

W. Li, G. Dasarathy, K. Natesan Ramamurthy, and V. Berisha. Finding the homology of decision boundaries with active learning. In H. Larochelle, M. Ranzato, R. Hadsell, M. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 8355–8365. Curran Associates, Inc., 2020. URL https://proceedings.neurips.cc/paper_files/paper/2020/file/5f14615696649541a025d3d0f8e0447f-Paper.pdf.

M. Liu, T. K. Dey, and D. F. Gleich. Topological structure of complex predictions. *Nature Machine Intelligence*, Nov 2023a. ISSN 2522-5839. doi: 10.1038/s42256-023-00749-8. URL https://doi.org/10.1038/s42256-023-00749-8.

Y. Liu, C. M. Cole, C. Peterson, and M. Kirby. ReLU neural networks, polyhedral decompositions, and persistent homology. In *2nd Annual TAG in Machine Learning*, 2023b.

Z. Liu, P. Luo, X. Wang, and X. Tang. Deep learning face attributes in the wild. In *Proceedings of International Conference on Computer Vision (ICCV)*, December 2015.

V. Lohweg. Banknote Authentication. UCI Machine Learning Repository, 2013. DOI: https://doi.org/10.24432/C55P57.

D. Loiseaux, M. Carrière, and A. Blumberg. A framework for fast and stable representations of multiparameter persistent homology decompositions. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023a. URL `https://openreview.net/forum?id=nKCUDd9GYu`.

D. Loiseaux, L. Scoccola, M. Carrière, M. B. Botnan, and S. Oudot. Stable vectorization of multiparameter persistent homology using signed barcodes as measures. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023b. URL `https://openreview.net/forum?id=4mwORQjAim`.

E. Love, B. Filippenko, V. Maroulas, and G. E. Carlsson. Topological convolutional neural networks. In *TDA & Beyond*, 2020. URL `https://openreview.net/forum?id=hntbh8Zo1V`.

D. Lütgehetmann, D. Govc, J. P. Smith, and R. Levi. Computing persistent homology of directed flag complexes. *Algorithms*, 13(1), 2020. ISSN 1999-4893. doi: 10.3390/a13010019. URL `https://www.mdpi.com/1999-4893/13/1/19`.

M. Masden. Algorithmic determination of the combinatorial structure of the linear regions of ReLU neural networks, 2022. URL `https://arxiv.org/abs/2207.07696`.

L. Matthey, I. Higgins, D. Hassabis, and A. Lerchner. dsprites: Disentanglement testing sprites dataset. https://github.com/deepmind/dsprites-dataset/, 2017.

G. Montufar, N. Otter, and Y. G. Wang. Can neural networks learn persistent homology features? In *TDA & Beyond*, 2020. URL `https://openreview.net/forum?id=pqpXM1Wjsxe`.

M. Muller, S. Kroon, and S. Chalup. Topological dynamics of functional neural network graphs during reinforcement learning. In B. Luo, L. Cheng, Z.-G. Wu, H. Li, and C. Li, editors, *Neural Information Processing*, pages 190–204, Singapore, 2024. Springer Nature Singapore. ISBN 978-981-99-8138-0.

J. R. Munkres. *Topology*. Prentice Hall, Inc., Upper Saddle River, NJ, 2nd edition, 2000. ISBN 0131816292; 9780131816299; 0131784498; 9780131784499.

G. Naitzat, A. Zhitnikov, and L.-H. Lim. Topology of deep neural networks. *J. Mach. Learn. Res.*, 21(1), Jan 2020. ISSN 1532-4435.