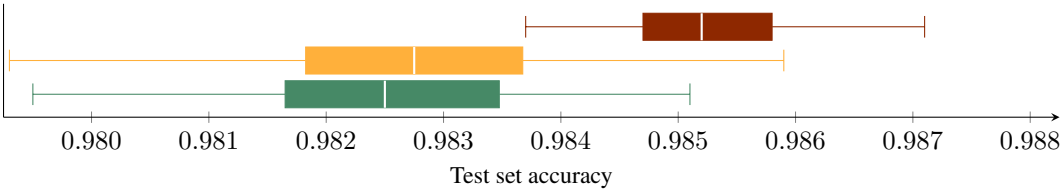

Figure A.13: Comparison of test set accuracy for trained networks without modifications (green), with batch normalization (yellow), and with 50% of the neurons dropped out during training (red) for the MNIST data set.

A.8 Testing accuracy of differently regularized models

We showed in the main text that neural persistence is capable of distinguishing between networks trained with/without batch normalization and/or dropout. Figure A.13 additionally shows test set accuracies.